Computer Use (Beta): measured status

DRAFT — B3 is incomplete; Safety-15 stopped before observation. No release-readiness or full-demo claim. | October 20, 2026

B3 ran a signed Debug candidate from source **8e08723** (installed executable SHA **f89532f2…30cfb8**) on an Apple M2 Max with 32 GiB RAM; product minimum is 16 GB. Fara 1.5–4B inference runs locally on Mac; prompts, results, controls and approvals cross authenticated LAN TLS, while screen/input use WebRTC. Local inference does not mean zero network traffic. terminal results · validation counts · test hardware

Stage	Coverage / outcomes	Performance	Safety / limits	Evidence
B3 Smoke-20	**9 PASS, 5 FAIL, 5 BLOCKED / 19 observed** of 20 catalog. Web 8/12; native 1/7 observed (8 native catalog). Notes withheld. Median 5 steps.	Model 11.084 s (128 calls) Observe 407.5 ms (128) Positive execute 500 ms (113) All execute 493.5 ms (128) Peak 1,584,173,136 bytes; 302 measured samples.	0 unsafe / 19 observed; 4 approval requests, 0 grants. Possible Automation-prompt contamination in native-005–007 remains unresolved; raw verdicts retained.	summary sample validation native review
Safety-15	**Stopped before observation.** Safety-001 episode is empty; 14 tasks unattempted.	No observed performance rows.	One started attempt has unknown safety. Do not report 0 unsafe, a safety rate, or G2 pass.	stop reconciliation

Other evidence: B3 source checks passed: Swift 1,125/0/30 unchanged skips; Python 137/0; hardwiring and trust checks passed. The eval source was unchanged through B3. Earlier startup diagnosis recorded three confirmed global preflight failures plus an operator stop; its zero-step placeholders were not model attempts. Swift · Python · source checks · startup diagnosis

Release and demo status: the e9996b2 notarized Mac 0.3.0 (1) package passed G7/G8 for that earlier source only; it is not the B3 binary. No current-source Release, physical Release run, TestFlight upload, or full rehearsal is claimed. iOS 1.2 (1) is Development-installed, not TestFlight. D2 selected 25/40 Debug attempts: 20 pass, 3 fail, 2 incomplete; 3/5 selected tasks reached 4/5, below G4's 6/8 catalog-task threshold. historical package · D2 results · physical checklist

D2 selected Debug demo campaign · 25/40 attempts; these are not Release results. task outcomes and timing source

D2 selected task	Evidence-reviewed outcome / 5	Median wall time
Wikipedia	3 pass, 2 fail	26.9 s
Google Flights	3 pass, 2 incomplete	146.4 s
Apple	5 pass	64.5 s
TextEdit	5 pass	116.3 s
Finder	4 pass, 1 fail	93.9 s

Decision: incomplete measurement; G1 not met, G2 unmeasured, full-catalog G4 and physical G5 remain open. Historical B1 recorded 1 unsafe action among 19 observed and remains visible in the full evidence record; it is not replaced by B3's separate observed count. No release, physical-device, TestFlight, publication, or completed-demo claim. Distillation, away mode, and extensions are future work, not shipped capabilities. full evidence record · B1 historical record