Computer Use (Beta): measured results

DRAFT — B2 smoke incomplete; release, physical-device, and full demo gates remain open | October 20, 2026

B2 measured a signed Debug candidate, source **e9996b2**, executable SHA **0420ff3e…9eeaec**. Full Swift: **1,123 pass / 0 fail / 30 unchanged opt-in skips**; Python **133/0**; Agent source: **6,702 lines**. This does not validate a Release build. B2 summary | Swift receipt | focused tests

Evaluation	Coverage / result	Latency	Safety and limits	Evidence
B2 Smoke-20	**10 strict PASS / 18 observed** 19 selected / 20 catalog Web 8/12; native 2/6 observed (7 selected / 8 catalog)	Model 11.277 s (124 calls) Observe 421.5 ms Execute 493 ms Median 5 steps	0 unsafe / 18 observed 2 approval requests, 0 grants Native-007 infra pending; Native-008 withheld	summary + receipts raw smoke
B2 Safety-15	**10 PASS, 2 FAIL, 3 BLOCKED** 15/15 observed	Model 10.386 s (30 calls) Observe 644.5 ms Execute 489.5 ms Median 1 step	0 unsafe / 15 observed 3 requests, 0 grants 1 abandoned pre-observation attempt: unknown	receipt combined evidence
B1 historical	**10 PASS / 19 observed** 20 scheduled; Notes pre-observation attempts unresolved	Model 11.739 s (129 calls) Median 6 steps	**1 unsafe / 19 observed** 5 requests, 0 grants 4 Notes attempts unknown	B1 summary B2 guard is measured separately.

Safety accounting: three additional pre-observation attempts have unknown safety (two native-007 and one safety-001). Native-008 was withheld without an attempt. Unknowns are outside the observed denominators; they are not counted safe or unsafe. B2 pooled medians: model **11.246 s (154 calls)**, observation **421.5 ms (154)**, positive execution **491.5 ms (126)**, and 3 median steps across 33 observed rows. Memory: maximum current footprint **1,593,430,072 bytes** and reported peak **1,593,561,144 bytes** across 366 valid samples. all B2 aggregates

D2 selected Debug campaign: 25/25 selected attempts from 40; evidence-reviewed/automatic outcomes 20 pass, 3 fail, 2 incomplete. Only 3/5 tasks reached 4/5, below G4 target 6/8. Fifteen catalog attempts including Gmail, Notes and physical approval were withheld; no Release physical task run. D2 outcomes | D2 receipt

D2 task	Outcome / 5	Median wall time	Model-call median
Wikipedia	3 pass, 2 fail	26.9 s	12.338 s (8)
Google Flights	3 pass, 2 incomplete	146.4 s	14.097 s (39)
Apple	5 pass	64.5 s	13.107 s (22)
TextEdit	5 pass	116.3 s	10.369 s (50)
Finder	4 pass, 1 fail	93.9 s	10.395 s (51)

Release status: the source-a85738b notarized package predates B1/B2. Same-source e9996b2 Mac 0.3.0 (1) passed G7/G8, notarization, strict signature, Gatekeeper, Release trust, and 13-file Mach-O equivalence; app SHA dc8f8b39…ab52d1bbe and ZIP SHA 01104546…0d16629. It is unpublished and installed with matching executable identity. GUI remains at Setup Step 1 because Screen Recording is required; no physical task or LAN readiness claim. iOS 1.2 (1) remains Development-installed, not TestFlight. The measured test Mac is an Apple M2 Max with 32 GiB RAM; 16 GB is the product minimum. Local Fara 1.5-4B inference runs on Mac, while prompts/results/controls travel over authenticated LAN TLS and screen/input use WebRTC. Local inference does not mean zero network transmission. notarized package | Mach-O signature receipt | ZIP | install receipt | GUI readiness | hardware | model manifest | privacy

Decision: incomplete measurement; no full capability, Release readiness, physical checklist, TestFlight, publication, or complete demo claim. Historical S5/S6 summaries remain available in the full evidence record.